

APPLE TWITTER SENTIMENT ANALYSIS

INTRODUCTION

In today's world, we are living in digital times. Social media platforms have been used as a source of public opinions and sentiments. Apple and Google companies monitor opinions and sentiments from customers and analyse the customer satisfaction so as to identify areas of improvement. These sentiments can offer good insights for business to improve their products and manage the brand image. This project aims to develop a Natural Language Processing (NLP) model to analyse Twitter sentiments about Apple and Google products.

Using dataset from CrowdFlower, I will train the model to predict the sentiment of any given tweet based on the content. This will help the Stakeholder to make data-driven decisions.

BACKGROUND

Sentiment analysis helps in identifying and extracting subjective information from a text. Companies use Sentiment analysis to analyse customer feedback and make data-driven decisions. This project uses tweets about Apple and Google products rated by humans for sentiments to train and evaluate the NLP model.

DATA OVERVIEW

The dataset was sourced from CrowdFlower via data.world, the link for Kaggle is <https://www.kaggle.com/datasets/slythe/apple-twitter-sentiment-crowdflower> (<https://www.kaggle.com/datasets/slythe/apple-twitter-sentiment-crowdflower>).

There are 3886 Records and 12 Features.

The key Features are :

1. sentiment - sentiment labels
2. sentiment:confidence - confidence score for sentiment
3. date- when tweet was posted
4. query- search query used
5. text- actual tweet

The sentiments are already labelled and classified into classes which include:

1. '1' representing 'Negative'
2. '3' representing 'Neutral'
3. '5' representing 'Positive'
4. 'not_relevant' representing 'not relevant'

BUSINESS PROBLEM

It is crucial for companies to understand customer sentiments and enhance their satisfaction. Apple and Google company is interested in understanding the customer sentiments so as to improve their products and enhance customer satisfaction. By building an NLP model that accurately predicts sentiments of the tweets can help the company to make informed decision which can improve their performance hence contributing to business growth.

OBJECTIVE

Build a model that can rate sentiment of a tweet based on the content.

SUMMARY

1. Business and Data Understanding

The project aims in sentiment analysis. It analyses Apple and Google products and classify them into Positive(5), Negative(1), Neutral(3), Non relevant(non_relevant). The data consist of labeled tweets with sentiments making it well suited for classification task. Descriptive analysis was used such as distribution of sentiment classes. The data sourced from Twitter enables the company to gauge the customer satisfaction and get an insight of areas to improve to enhance the company performance.

2. Data Preparation

- The process included Data understanding, Data Cleaning and Text processing.
- Data understanding helped to know the size of dataset i was working on, the data structures i.e shape and identify the data types and whether its a categorical or a numerical feature.
- The data contained 3886 Records and 12 Features. The datatypes were bool(1), float64(2), int64(2), object(7).
- The dataset had 2 columns which had missing data and also the columns were irrelevant according to my objective so i dropped the columns plus other columns which were relevant and remained with two columns which were sentiment and text.
- Text preprocessing to remove URLs, Hashtags, mentions, special characters which helped to reduce noise and focus on the actual content of the tweet and converting text to lowercase to reduce vocabulary size and improve word frequency analysis.
- Stopwords were removed and Lemmatization together with stemming applied to reduce word to their base form.
- The library used for prepare the data were NLTK for Tokenization, stopwords removal, Lemmatization, Stemming and TfidfVectorizer from Scikit learn to convert text to Numeric.

3. Modeling

I used two modeling packages which included Scikit-learn and Natural Language Toolkit(NLTK). Sklearn models included:

- Logistic Regression which is simple and interpretable and effective for text classification task,I used IT in Binary classification which then expanded to multiclass classification.
- Decision tree and RandomForest was performed and ensemble methods like Bagging to improve the prediction accuracy though it did not help improve the accuracy.it helped reduce overfitting.
- XGBoost was used to capture complex patterns but did not improve the model over Logistic Regression.
- NLTK was used for tokenization,Lemmatization,stemming and Stopwords removal.
- Hyperparameter Tuning was performed using GridsearchCV to find optimal parameters.Performed cross validation with 5 folds to ensure model generalized well on unseen data.
- Feature engineering with Tfidfvectorizer with different n_gram range.
- Removed stop words to reduce noise and improve performance.
- Logistic Regression with class_weight ='balanced' was used to handle the class imbalance and improve performance.

4. Evaluation

- The models were evaluated with Accuracy,Precision,Recall and F1 score.
- The best model was Logistic Regression with Gridsearch,it had Accuracy of 72% ,precision of 68% and Bagged RandomForestClassifier with Accuracy 71% and precision 70%.
- In the Validation approach,I used a Train_test split to validate the model performance,ensuring the results are well generalised with new data.The dataset was split into 70% training data and 30% test data.

Data Understanding

Data Preparation

This process involve preparing the data which include Loading the dataset,checking the data types,checking the shape,calculating the summary statistics.

```
In [1]: ▶ # Install XGboost
        from xgboost import XGBClassifier
```

Data Description

```
In [2]: # Import relevant columns  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
from sklearn.preprocessing import LabelEncoder,OneHotEncoder,StandardScaler  
from sklearn.metrics import accuracy_score,recall_score,precision_score,f1_  
from sklearn.model_selection import train_test_split,GridSearchCV  
from sklearn.ensemble import RandomForestClassifier,BaggingClassifier  
from sklearn.pipeline import Pipeline  
import re  
import nltk  
from nltk.corpus import stopwords  
from nltk.tokenize import RegexpTokenizer  
from nltk.stem import WordNetLemmatizer,PorterStemmer  
from sklearn.feature_extraction.text import TfidfVectorizer,CountVectorizer  
from sklearn.linear_model import LogisticRegression  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.naive_bayes import MultinomialNB  
nltk.download('stopwords')  
nltk.download('wordnet')
```

```
[nltk_data] Downloading package stopwords to  
[nltk_data]   C:\Users\user\AppData\Roaming\nltk_data...  
[nltk_data]   Package stopwords is already up-to-date!  
[nltk_data] Downloading package wordnet to  
[nltk_data]   C:\Users\user\AppData\Roaming\nltk_data...  
[nltk_data]   Package wordnet is already up-to-date!
```

Out[2]: True

```
In [3]: # Loading the dataset
apple_df = pd.read_csv("Apple-Twitter-Sentiment-DFE.csv",encoding='latin1')
apple_df.head() # checking the first 5 rows
```

```
Out[3]:
```

	_unit_id	_golden	_unit_state	_trusted_judgments	_last_judgment_at	sentiment	sentiment:confidence
--	----------	---------	-------------	--------------------	-------------------	-----------	----------------------

0	623495513	True	golden	10	NaN	3	0.99
---	-----------	------	--------	----	-----	---	------

1	623495514	True	golden	12	NaN	3	0.99
---	-----------	------	--------	----	-----	---	------

2	623495515	True	golden	10	NaN	3	0.99
---	-----------	------	--------	----	-----	---	------

3	623495516	True	golden	17	NaN	3	0.99
---	-----------	------	--------	----	-----	---	------

4	623495517	False	finalized	3	12/12/14 12:14	3	0.99
---	-----------	-------	-----------	---	----------------	---	------



```
In [4]: # Checking the columns
apple_df.columns
```

```
Out[4]: Index(['_unit_id', '_golden', '_unit_state', '_trusted_judgments',
              '_last_judgment_at', 'sentiment', 'sentiment:confidence', 'date',
              'id',
              'query', 'sentiment_gold', 'text'],
              dtype='object')
```

In [5]: `# Checking information`
`apple_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3886 entries, 0 to 3885
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   _unit_id                             3886 non-null   int64
1   _golden                             3886 non-null   bool
2   _unit_state                         3886 non-null   object
3   _trusted_judgments                  3886 non-null   int64
4   _last_judgment_at                  3783 non-null   object
5   sentiment                           3886 non-null   object
6   sentiment:confidence               3886 non-null   float64
7   date                               3886 non-null   object
8   id                                 3886 non-null   float64
9   query                              3886 non-null   object
10  sentiment_gold                     103 non-null    object
11  text                               3886 non-null   object
dtypes: bool(1), float64(2), int64(2), object(7)
memory usage: 337.9+ KB
```

In [6]: `# Checking the shape`
`apple_df.shape`

Out[6]: (3886, 12)

- The dataset has 3886 rows and 12 columns.
- The data types are 1 boolean, 2 float64, 2 int64 and 7 object.
- The dataset has columns that has missing values like `_last_judgment_at`, `sentiment_gold`
- There is also unnecessary columns for our use case.

In [7]: `# Calculating the summary statistics`
`apple_df.describe()`

Out[7]:

	<code>_unit_id</code>	<code>_trusted_judgments</code>	<code>sentiment:confidence</code>	<code>id</code>
count	3.886000e+03	3886.000000	3886.000000	3.886000e+03
mean	6.234975e+08	3.687082	0.829526	5.410039e+17
std	1.171906e+03	2.004595	0.175864	7.942752e+14
min	6.234955e+08	3.000000	0.332700	5.400000e+17
25%	6.234965e+08	3.000000	0.674475	5.400000e+17
50%	6.234975e+08	3.000000	0.811250	5.410000e+17
75%	6.234984e+08	3.000000	1.000000	5.420000e+17
max	6.235173e+08	27.000000	1.000000	5.420000e+17

```
In [8]: ▶ # Checking the sentiment classes
print(apple_df['sentiment'].unique())

['3' '5' '1' 'not_relevant']
```

The sentiment classes include:

- '1' representing 'Negative'
- '3' representing 'Neutral'
- '5' representing 'Positive'
- 'not_relevant' representing 'not relevant'


```
In [9]: ▶ #Checking the distribution of the classes
counts = apple_df['sentiment'].value_counts()
counts
```

```
Out[9]: sentiment
3          2162
1          1219
5           423
not_relevant    82
Name: count, dtype: int64
```

```
In [10]: # Visualise the sentiment classes
labels = counts.index
values = counts.values


labels = ['Neutral (3)', 'Negative (1)', 'Positive (5)', 'Not Relevant']

plt.figure(figsize=(10,6))
plt.bar(labels,counts.values,color = ['g','b','r','gray'])
plt.title('Distribution of Sentiment classes')
plt.xlabel('sentiment classes')
plt.ylabel('count')
plt.show()
```



- Neutral(3) had the highest sentiment, followed by Negative sentiment(1), then Positive(5) and the non_relevant which had the smallest.


```
In [11]: # View the distribution of sentiment confidence
sns.histplot(apple_df['sentiment:confidence'], bins=30)
plt.title('Distribution of Sentiment Confidence Scores')
plt.show()
```



- The sentiment confidence of 1 has the highest frequency scores indicating there is many instances where the model is highly confident about the sentiment prediction.
- There is also an indication of moderate level of confidence between 0.5 to 0.8 and also smaller peaks around 0.3 to 0.4 indicating low to moderate confidence.

Data Cleaning

Missing values

```
In [12]: # Checking the missing values
apple_df.isna().sum()
```

```
Out[12]: _unit_id          0
         _golden          0
         _unit_state      0
         _trusted_judgments 0
         _last_judgment_at 103
         sentiment        0
         sentiment:confidence 0
         date             0
         id               0
         query            0
         sentiment_gold    3783
         text             0
         dtype: int64
```

```
In [13]: # Drop the missing values and irrelevant columns
apple_df= apple_df.drop(columns=['sentiment_gold','_last_judgment_at'],axis=1)
```

```
In [14]: # Drop irrelevant columns
apple_cleaned_df = apple_df.drop(columns=['_unit_id', '_golden', '_unit_status'])
apple_cleaned_df.dropna().head() # Drop missing rows and display the first
```

Out[14]:

	sentiment	text
0	3	#AAPL:The 10 best Steve Jobs emails ever...htt...
1	3	RT @JPDesloges: Why AAPL Stock Had a Mini-Flas...
2	3	My cat only chews @apple cords. Such an #Apple...
3	3	I agree with @jimcramer that the #IndividualIn...
4	3	Nobody expects the Spanish Inquisition #AAPL

Duplicated columns

```
In [15]: # Drop the duplicates
apple_cleaned_df= apple_cleaned_df.drop_duplicates()
```

Text Processing

```
In [16]: # Feature Engineering
# Tokenize
tokenizer = RegexpTokenizer(r'\w+')
# Create a list of stopwords in English
stopwords_list = stopwords.words('English')
# Initialize Lemmatizer
lemmatizer = WordNetLemmatizer()
```

```
In [17]: ▶ def preprocess_text(text, tokenizer, stopwords_list, lemmatizer):
# Remove URLs
text = re.sub(r"http\S+|www\S+", "", text)
# Remove hashtags
text = re.sub(r"#\S+", " ", text)
# Remove mentions
text = re.sub(r"@S+", " ", text)
# Remove special characters and numbers
text = re.sub(r"[^a-zA-Z\s]", " ", text)
# Convert to lowercase
text = text.lower()
# Tokenize
tokens = tokenizer.tokenize(text)
# Remove stop words
tokens = [word for word in tokens if word not in stopwords_list]
# Lemmatize
tokens = [lemmatizer.lemmatize(word) for word in tokens]
# Join tokens back to string
text = " ".join(tokens)
return text
```

```
In [18]: ▶ # Apply the preprocessing function to the DataFrame
apple_cleaned_df['prepro_text'] = apple_cleaned_df['text'].apply(lambda x:
apple_cleaned_df.head()
```

```
Out[18]:
```

	sentiment	text	prepro_text
0	3	#AAPL:The 10 best Steve Jobs emails ever...htt...	best steve job email ever
1	3	RT @JPDesloges: Why AAPL Stock Had a Mini-Flas...	rt aapl stock mini flash crash today aapl
2	3	My cat only chews @apple cords. Such an #Apple...	cat chew cord
3	3	I agree with @jimcramer that the #IndividualIn...	agree trade extended today pullback good see
4	3	Nobody expects the Spanish Inquisition #AAPL	nobody expects spanish inquisition

Visualize a sample text

```
In [19]:  for i in range (10):
            print(f"original text: {apple_cleaned_df['text'].iloc[i]}")
            print(f"cleaned text: {apple_cleaned_df['prepro_text'].iloc[i]}")
            print(" ")
```

original text: #AAPL:The 10 best Steve Jobs emails ever...<http://t.co/82G1kl94tx>
cleaned text: best steve job email ever

original text: RT @JPDesloges: Why AAPL Stock Had a Mini-Flash Crash Today \$AAPL #aapl
<http://t.co/hGFcjYa0E9> (<http://t.co/hGFcjYa0E9>)
cleaned text: rt aapl stock mini flash crash today aapl

original text: My cat only chews @apple cords. Such an #AppleSnob.
cleaned text: cat chew cord

original text: I agree with @jimcramer that the #IndividualInvestor should not trade #Apple #AAPL, it's extended so today's pullback is good to see
cleaned text: agree trade extended today pullback good see

original text: Nobody expects the Spanish Inquisition #AAPL
cleaned text: nobody expects spanish inquisition

original text: #AAPL:5 Rocket Stocks to Buy for December Gains: Apple and More...<http://t.co/eG5XhXdLLS>
cleaned text: rocket stock buy december gain apple

original text: Top 3 all @Apple #tablets. Damn right! <http://t.co/RJiGn2JUuB> (<http://t.co/RJiGn2JUuB>)
cleaned text: top damn right

original text: CNBCTV: #Apple's margins better than expected? #aapl <http://t.co/7geVrtOGLK> (<http://t.co/7geVrtOGLK>)
cleaned text: cnbctv margin better expected

original text: Apple Inc. Flash Crash: What You Need to Know <http://t.co/YJlgtifdAj> (<http://t.co/YJlgtifdAj>) #AAPL
cleaned text: apple inc flash crash need know

original text: #AAPL:This Presentation Shows What Makes The World's Biggest Tech Companies ...<http://t.co/qlH9PqSoSd>
cleaned text: presentation show make world biggest tech company

Building a model

I will first limit my analysis to Positive and Negative tweets only building a binary classifier and then add the Neutral tweets to build a Multiclass classifier.

```
In [20]: print(apple_cleaned_df['sentiment'].unique())

['3' '5' '1' 'not_relevant']
```

Binary class classifier

```
In [21]: # Filter to include positive(5) and Negative(1) sentiments only
binary_df = apple_cleaned_df[apple_cleaned_df['sentiment'].isin(['5','1'])]
binary_df.head()
```

```
Out[21]:
```

	sentiment	text	prepro_text
6	5	Top 3 all @Apple #tablets. Damn right! http://...	top damn right
7	5	CNBCTV: #Apple's margins better than expected?...	cnbctv margin better expected
10	1	WTF MY BATTERY WAS 31% ONE SECOND AGO AND NOW ...	wtf battery one second ago wtf
13	5	RT @peterpham: Bought my @AugustSmartLock at t...	rt bought store pretty good logo match wait in...
14	1	@apple Contact sync between Yosemite and iOS8 ...	contact sync yosemite io seriously screwed use...

Define X and y

```
In [22]: X_bin = binary_df['prepro_text']
X_bin.head()
```

```
Out[22]: 6          top damn right
7          cnbctv margin better expected
10         wtf battery one second ago wtf
13  rt bought store pretty good logo match wait in...
14  contact sync yosemite io seriously screwed use...
Name: prepro_text, dtype: object
```

```
In [23]: y_bin = binary_df['sentiment']
```

Pipeline

- Create a pipeline for vectorizing and also Hyper tuning
- used class_weight='balanced' to balance the classes

```
In [24]: # Define pipeline
pipe = Pipeline([
    ('vectorizer', TfidfVectorizer(ngram_range=(1,2))),
    ('model', LogisticRegression(random_state=42, class_weight='balanced')),
])
# Define parameter grid
param_grid = [
    {
        'model': [LogisticRegression(max_iter=200, solver='liblinear', class_weight='balanced')],
        'model__C': [0.001, 0.01, 0.1, 1, 10, 100, 1000]
    },
    {
        'model': [RandomForestClassifier(class_weight='balanced')],
        'model__n_estimators': [50, 100],
        'model__criterion': ['gini', 'entropy'],
        'model__max_depth': [None, 10, 20]
    },
    {
        'model': [DecisionTreeClassifier(random_state=42)],
        'model__criterion': ['gini', 'entropy'],
        'model__max_depth': [None, 10, 20]
    }
]
```

Train-Test split

```
In [25]: # Split the data into training and test data
X_train_bin, X_test_bin, y_train_bin, y_test_bin = train_test_split(X_bin, y_bin,
```

Model Training and fitting(binary)

```
In [26]: pipe.fit(X_train_bin, y_train_bin)
```

```
Out[26]: Pipeline(steps=[('vectorizer', TfidfVectorizer(ngram_range=(1, 2))),
                          ('model',
                           LogisticRegression(class_weight='balanced', random_state=42))])
```

In [27]: `print(X_train_bin)`

```
1131      hour hold customer service trying figure messup
1355      lol fuck going pay left arm amp leg product ma...
2657      ok look yall still reversed back io idk hold i...
2073      leave mac plugged leave house need check power...
1122      phone die plugged wtf
...
2859      freaking computer went psycho ruined life than...
3241      day without access passwd apple working new ap...
2343      duck
3816      rt phone went please tell mathematically possible
2847      company official acct company
Name: prepro_text, Length: 1036, dtype: object
```

In [28]: `# Predict`
`y_pred_binary = pipe.predict(X_test_bin)`

Model Evaluation

In [29]: `# Convert target labels to integers`
`y_test_bin = y_test_bin.astype(int)`
`y_pred_binary = y_pred_binary.astype(int)`

In [30]: `accuracy_bin = accuracy_score(y_test_bin,y_pred_binary)`
`precisin_bin = precision_score(y_test_bin,y_pred_binary)`
`recall_bin = recall_score(y_test_bin,y_pred_binary)`
`f1_bin = f1_score(y_test_bin,y_pred_binary)`
`roc_auc_bin = roc_auc_score(y_test_bin,y_pred_binary)`

`print(f"Accuracy : {accuracy_bin}")`
`print(f"Precision : {precisin_bin}")`
`print(f"Recall : {recall_bin}")`
`print(f"F1 : {f1_bin}")`
`print(f"ROC_AUC : {roc_auc_bin}")`

```
Accuracy : 0.7955056179775281
Precision : 0.9006211180124224
Recall : 0.830945558739255
F1 : 0.8643815201192251
ROC_AUC : 0.7488061127029607
```

In [31]:


```
classif_report = classification_report(y_test_bin,y_pred_binary)
conf_matrix = confusion_matrix(y_test_bin,y_pred_binary)

print(f"classification_Report:,{classif_report}")
print(f"confusion_matrix:",conf_matrix)
```

classification_Report: ,			precision	recall	f1-score	supp
ort						
	1	0.90	0.83	0.86		349
	5	0.52	0.67	0.58		96
	accuracy			0.80		445
	macro avg	0.71	0.75	0.72		445
	weighted avg	0.82	0.80	0.80		445

confusion_matrix: [[290 59]
[32 64]]


```
In [32]: ▶ # Plot confusion matrix
plt.figure(figsize=(12,10))
sns.heatmap(conf_matrix,annot=True,cmap="cividis",xticklabels=['5', '1'],yt
plt.xlabel('predicted label')
plt.ylabel('actual label')
plt.show()
```



Interpretation:

- The score had accuracy of 80%
- The model performs well for the negative class (1) with high precision, recall, accuracy, and F1 score.
- Performance for the positive class (5) is lower, indicating some challenges in correctly identifying positive tweets.

Join to form a Multi class classifier

```
In [33]: # Create a dataframe containing the four classes of sentiment
multi_df = apple_cleaned_df[apple_cleaned_df['sentiment'].isin(['3', '5', '1'])]
multi_df.head()
```

```
Out[33]:
```

	sentiment	text	prepro_text
0	3	#AAPL:The 10 best Steve Jobs emails ever...htt...	best steve job email ever
1	3	RT @JPDesloges: Why AAPL Stock Had a Mini-Flas...	rt aapl stock mini flash crash today aapl
2	3	My cat only chews @apple cords. Such an #Apple...	cat chew cord
3	3	I agree with @jimcramer that the #Individualln...	agree trade extended today pullback good see
4	3	Nobody expects the Spanish Inquisition #AAPL	nobody expects spanish inquisition

```
In [34]: # Filter out the non relevant class
multi_df = multi_df[multi_df['sentiment'] != 'not_relevant']
```

- Filtered out the non_relevant class as it was not relevant to the Apple and Google products so it was bringing noise to the data hence reducing the modeling performance

```
In [35]: print(multi_df['sentiment'].unique())

['3' '5' '1']
```

```
In [36]: X_multi = multi_df['prepro_text']
y_multi = multi_df['sentiment']
```

Train_Test Split

```
In [37]: X_train_mult,X_test_mult,y_train_mult,y_test_mult = train_test_split(X_multi,
X_train_mult)
```

```
Out[37]: 380      ikr u dont need bc hopefully buy one christmas...
1288      star analyst alex gauna jmp security rated aap...
2990      take convert shopper buyer amp create total se...
815              new patent broken screen
1303              domt think
...
3807              pick number one app year
1190              happy thnx fag
1225      apple poised record breaking holiday season aapl
1400      io fucked battery life iphone really starting ...
941              really
Name: prepro_text, Length: 2213, dtype: object
```

Model training and fitting(multi) With Gridsearch

```
In [38]: Pipe2 = Pipeline([
    ('vectorizer', TfidfVectorizer(ngram_range=(1,2))),
    ('model', LogisticRegression(random_state=42, class_weight='balanced'))
])
# Define parameter grid
param_grid = [
    {
        'model': [LogisticRegression(max_iter=200, solver='liblinear', multi_c
        'model__C': [1,10,100]
    },
    {
        'model' : [RandomForestClassifier(class_weight='balanced')],
        'model__n_estimators': [50,100],
        'model__criterion': ['gini', 'entropy'],
        'model__max_depth': [None,10,20]
    },
    {
        'model': [MultinomialNB()]
    }
]
# Perform Gridsearch
grid_search = GridSearchCV(Pipe2, param_grid, cv=5, verbose=1, n_jobs=-1, scoring=
grid_search.fit(X_train_mult, y_train_mult)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```

Out[38]: GridSearchCV(cv=5,
                      estimator=Pipeline(steps=[('vectorizer',
                                                  TfidfVectorizer(ngram_range=(1,
2))),
                                                  ('model',
                                                  LogisticRegression(class_weight
='balanced',
                                                                    random_state=4
2))]),
                      n_jobs=-1,
                      param_grid=[{'model': [LogisticRegression(C=1,
                                                                    class_weight='bala
nced',
                                                                    max_iter=200,
                                                                    multi_class='ovr',
                                                                    solver='liblinea
r')],
                                   'model__C': [1, 10, 100]},
                                   {'model': [RandomForestClassifier(class_weight
='balanced')],
                                   'model__criterion': ['gini', 'entropy'],
                                   'model__max_depth': [None, 10, 20],
                                   'model__n_estimators': [50, 100]},
                                   {'model': [MultinomialNB()]}],
                      scoring='accuracy', verbose=1)

```

Best Model Evaluation

```

In [39]: ▶ # Get the best model
          best_model = grid_search.best_estimator_
          best_model

```

```

Out[39]: Pipeline(steps=[('vectorizer', TfidfVectorizer(ngram_range=(1, 2))),
                          ('model',
                          LogisticRegression(C=1, class_weight='balanced', max_ite
r=200,
                                              multi_class='ovr', solver='liblinea
r'))])

```

```

In [40]: ▶ # Refit the best pipe to train
          best_model.fit(X_train_mult,y_train_mult)
          # predict the model
          y_pred_multi= best_model.predict(X_test_mult)
          accuracy_multi = accuracy_score(y_test_mult,y_pred_multi)

```

```
In [41]: accuracy_multi = accuracy_score(y_test_multi,y_pred_multi)
precision_multi = precision_score(y_test_multi,y_pred_multi,average='macro')
recall_multi = recall_score(y_test_multi,y_pred_multi,average='macro')
f1_multi = f1_score(y_test_multi,y_pred_multi,average='macro')

print(f"Accuracy : {accuracy_multi}")
print(f"Precision : {precision_multi}")
print(f"Recall : {recall_multi}")
print(f"F1 : {f1_multi}")
```

```
Accuracy : 0.7133825079030558
Precision : 0.6772349574362998
Recall : 0.6052073306600754
F1 : 0.6252947061062623
```

```
In [42]: classification_report_multi = classification_report(y_test_multi,y_pred_multi)
confusion_matrix_multi = confusion_matrix(y_test_multi,y_pred_multi)

print(f"classification Report:,{classification_report_multi}")
print(f"confusion matrix:",confusion_matrix_multi)
```


```
classification Report:,          precision    recall  f1-score   supp
ort

      1      0.71      0.66      0.68      312
      3      0.73      0.84      0.78      513
      5      0.59      0.31      0.41      124

 accuracy          0.71      949
 macro avg      0.68      0.61      0.63      949
weighted avg      0.70      0.71      0.70      949

confusion matrix: [[205 100   7]
 [ 60 433  20]
 [ 22  63  39]]
```

```
In [43]: ▶ # Plot confusion matrix
plt.figure(figsize=(12,10))
sns.heatmap(confusion_matrix_multi,annot=True,cmap="Blues",xticklabels=['3',
plt.xlabel('predicted label')
plt.ylabel('actual label')
plt.show()
```



Interpretation

- The accuracy of 71% indicates 71% of the predictions made by the model are correct.
- The precision of 68% indicates that, on average, 68% of the predicted positive instances are correct.
- The recall of 61% indicates on average, 63% of the actual positive instances are correctly identified by the model.
- The F1 score of 63% is a balance between precision and recall, providing a single metric that accounts for both false positives and false negatives.

Bagging(Bootstrap aggregation) with Decision Tree Classifier

```
In [44]: #convert text data(X_train_mult) to TF-IDF features  
vectorizer =TfidfVectorizer()  
X_train_tdif =vectorizer.fit_transform(X_train_mult)  
X_test_tdif =vectorizer.transform(X_test_mult)
```

```
In [45]: # Define the base model  
base_model = DecisionTreeClassifier()  
  
# Define the bagging ensemble model  
bagging_model = BaggingClassifier(base_model, n_estimators=50, random_state=42)  
  
# Train the model  
bagging_model.fit(X_train_tdif, y_train_mult)
```

```
Out[45]: BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=50,  
                           random_state=42)
```

```
In [46]: ▶ # Make predictions on the test set
y_pred_dt = bagging_model.predict(X_test_tdif)

# Calculate performance metrics
accuracy_dt = accuracy_score(y_test_mult, y_pred_dt)
precision_dt = precision_score(y_test_mult, y_pred_dt, average='macro')
recall_dt = recall_score(y_test_mult, y_pred_dt, average='macro')
f1_dt = f1_score(y_test_mult, y_pred_dt, average='macro')

print(f"Bagging Model_dt Accuracy: {accuracy_dt}")
print(f"Bagging Model_dt Precision: {precision_dt}")
print(f"Bagging Model_dt Recall: {recall_dt}")
print(f"Bagging Model_dt F1: {f1_dt}")

# Classification report and confusion matrix
classif_report_bagging_dt = classification_report(y_test_mult, y_pred_dt)
conf_matrix_bagging_dt = confusion_matrix(y_test_mult, y_pred_dt)

print(f"Bagging Model Classification Report_dt:{classif_report_bagging_dt}")
print(f"Bagging Model Confusion Matrix_dt:{conf_matrix_bagging_dt}")
```

Bagging Model_dt Accuracy: 0.6923076923076923

Bagging Model_dt Precision: 0.6401498320634252

Bagging Model_dt Recall: 0.5764688004359765

Bagging Model_dt F1: 0.5940905041748165

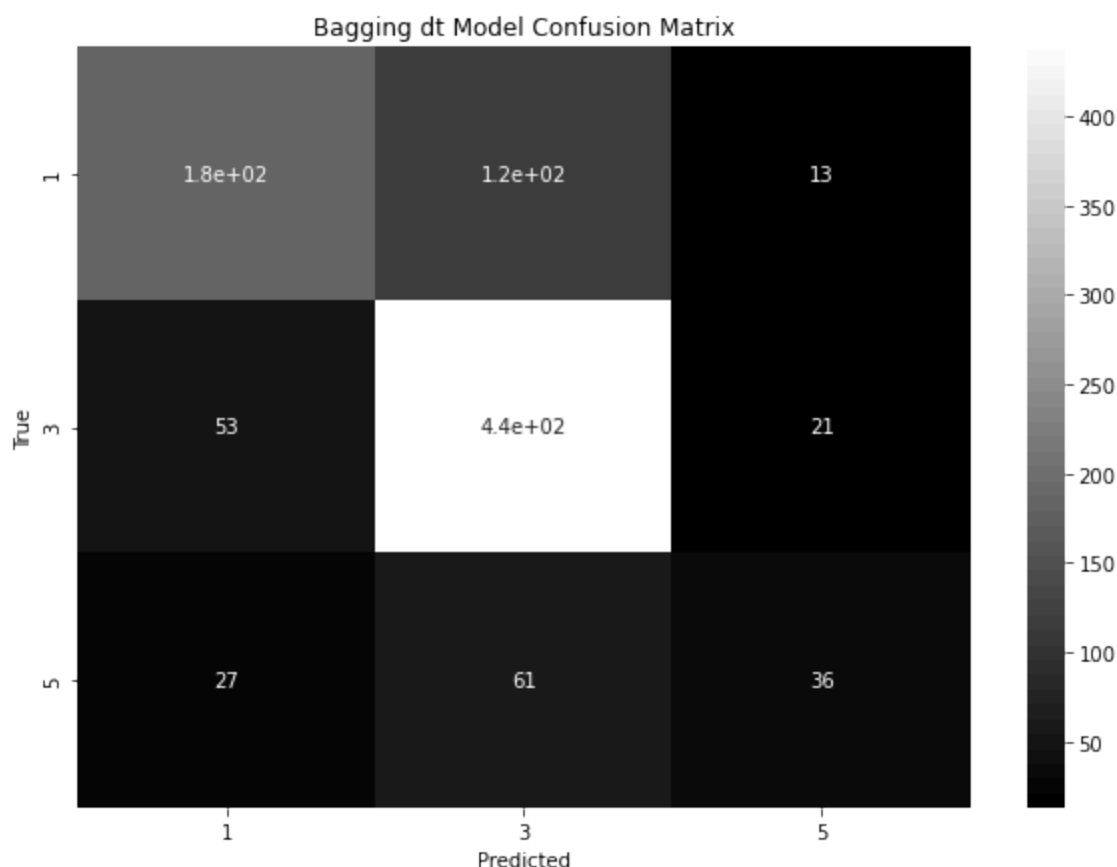
Bagging Model Classification Report_dt:					precision	recall
f1-score	support					

1	0.69	0.58	0.63	312
3	0.71	0.86	0.78	513
5	0.51	0.29	0.37	124

accuracy			0.69	949
macro avg	0.64	0.58	0.59	949
weighted avg	0.68	0.69	0.68	949

Bagging Model Confusion Matrix_dt:[[182 117 13]
[53 439 21]
[27 61 36]]


```
In [47]: ▶ plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix_bagging_dt, annot=True, cmap='gray', xticklabels=[
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Bagging dt Model Confusion Matrix')
plt.show()
```



Bagging with Random Forest Classifier

```
In [48]: ▶ rf = RandomForestClassifier(random_state=42)
```

```
In [49]: ▶ # Instantiate the RandomForestClassifier
RandomForestClassifier(n_estimators=200,bootstrap=True, min_samples_leaf=1,
max_depth=20, max_features='auto', random_state=42)
```

```
Out[49]: RandomForestClassifier(class_weight='balanced', max_depth=20, n_estimator
s=200,
random_state=42)
```

```
In [50]: ▶ # Fit the classifier
rf.fit(X_train_tdif,y_train_mult)
```

```
Out[50]: RandomForestClassifier(random_state=42)
```

```
In [51]: y_pred_rf = rf.predict(X_test_tdif)
```

```
In [52]: accuracy_rf = accuracy_score(y_test_mult, y_pred_rf)
precision_rf = precision_score(y_test_mult, y_pred_rf, average='macro')
recall_rf = recall_score(y_test_mult, y_pred_rf, average='macro')
f1_rf = f1_score(y_test_mult, y_pred_rf, average='macro')

print(f"Bagging Model_rf Accuracy: {accuracy_rf}")
print(f"Bagging Model_rf Precision: {precision_rf}")
print(f"Bagging Model_rf Recall: {recall_rf}")
print(f"Bagging Model_rf F1: {f1_rf}")

# Classification report and confusion matrix
classif_report_rf_bagging = classification_report(y_test_mult, y_pred_rf)
conf_matrix_rf_bagging = confusion_matrix(y_test_mult, y_pred_rf)

print(f"Bagging Model Classification Report_rf:{classif_report_rf_bagging}")
print(f"Bagging Model Confusion Matrix_rf:{conf_matrix_rf_bagging}")
```

```
Bagging Model_rf Accuracy: 0.708113804004215
Bagging Model_rf Precision: 0.6850221788781578
Bagging Model_rf Recall: 0.5674511904704321
Bagging Model_rf F1: 0.5880966047953209
```

```
Bagging Model Classification Report_rf:                precision      recall
f1-score      support
```

1	0.76	0.58	0.66	312
3	0.70	0.90	0.79	513
5	0.60	0.22	0.32	124
accuracy			0.71	949
macro avg	0.69	0.57	0.59	949
weighted avg	0.70	0.71	0.68	949

```
Bagging Model Confusion Matrix_rf:[[181 125   6]
[ 37 464  12]
[ 21  76  27]]
```

```
In [53]: ▶ # Confusion matrix visualization
plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix_rf_bagging, annot=True, cmap='magma', xticklabels=|
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Bagging rf Model Confusion Matrix')
plt.show()
```



XGBoost

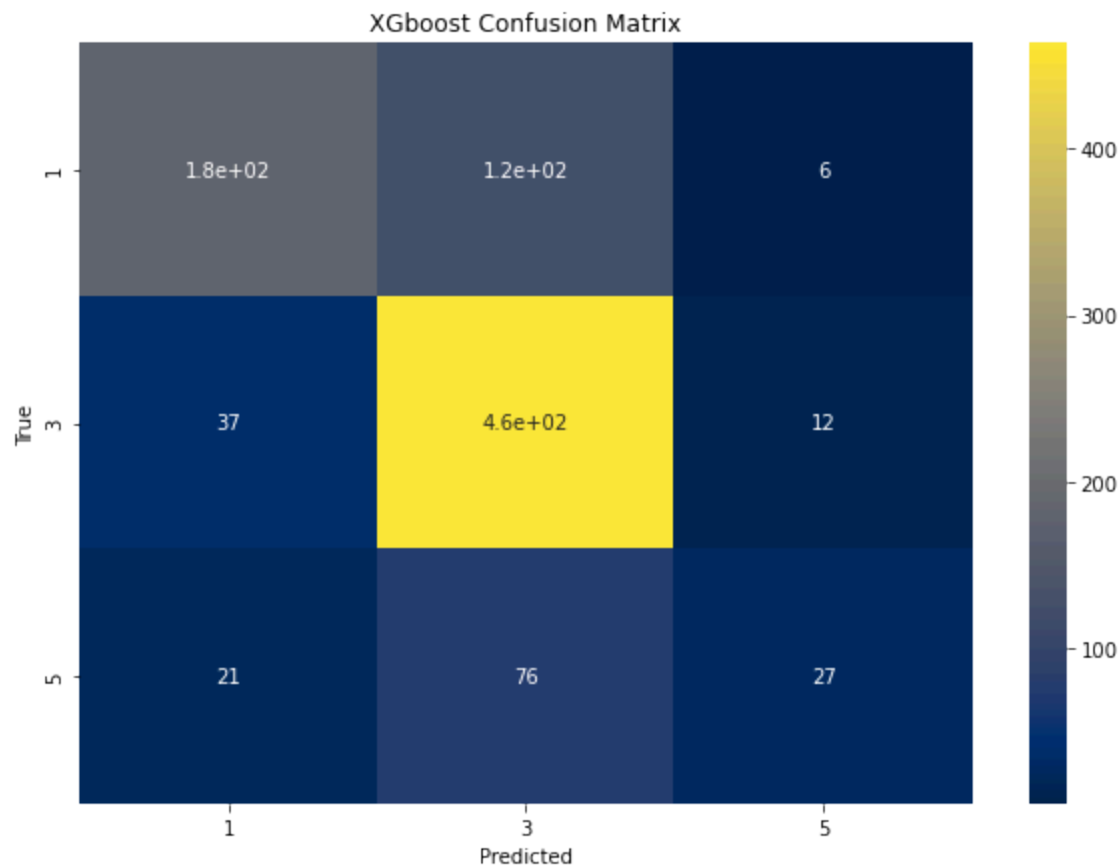
```
In [54]: ▶ # Create a pipeline
pipe3 = Pipeline([
    ('tfidf', TfidfVectorizer()),
    ('xgb', XGBClassifier(random_state =42))
])
# Initialize the Labelencoder
le = LabelEncoder()
# Fit and transform the pipeline
y_train_mult_encoded = le.fit_transform(y_train_mult)
y_test_mult_encoded = le.transform(y_test_mult)
# fit pipeline
pipe3.fit(X_train_mult, y_train_mult_encoded)
# Evaluate
y_pred_encoded = pipe3.predict(X_test_mult)
y_pred = le.inverse_transform(y_pred_encoded)
classif_report_XGboost = classification_report(y_test_mult, y_pred)
conf_matrix_XGboost = (confusion_matrix(y_test_mult, y_pred_rf))

print(f"Classification Report_XG:{classif_report_XGboost}")
print(f"Confusion Matrix_XG:{conf_matrix_XGboost}")
```

Classification Report_XG:				precision	recall	f1-score	su
pport							
	1	0.72	0.54	0.62		312	
	3	0.69	0.88	0.77		513	
	5	0.50	0.22	0.30		124	
	accuracy			0.68		949	
	macro avg	0.63	0.55	0.56		949	
	weighted avg	0.67	0.68	0.66		949	

Confusion Matrix_XG:[181 125 6]			
[37 464 12]			
[21 76 27]]			

```
In [55]: ▶ plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix_XGboost, annot=True, cmap='cividis', xticklabels=[
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('XGboost Confusion Matrix')
plt.show()
```



```
In [56]: ▶ # Create a dataframe to compare performance of the model
data = {
    'model': ['LogisticRegression', 'Bagging Ensemble(rf)', 'Bagging Ensemble
    'Accuracy': [71, 71, 69, 68],
    'Precision': [68, 69, 64, 63],
    'Recall': [61, 57, 58, 55],
    'F1': [63, 59, 59, 56]
}
df = pd.DataFrame(data)
df
```

Out[56]:

	model	Accuracy	Precision	Recall	F1
0	LogisticRegression	71	68	61	63
1	Bagging Ensemble(rf)	71	69	57	59
2	Bagging Ensemble(DT)	69	64	58	59
3	XGBoost	68	63	55	56

Save and Deploy a model

```
In [57]: ▶ import joblib

# # Save the trained model
model_name = 'Sentiment_model_joblib.pkl'
joblib.dump(best_model,model_name)

# # Save the vectorizer
vec_name = 'TfidfVectorizer_joblib.pkl'
joblib.dump(vectorizer,vec_name)
```

Out[57]: ['TfidfVectorizer_joblib.pkl']

```
In [58]: ▶ loaded_model = joblib.load('Sentiment_model_joblib.pkl')
loaded_vec = joblib.load('TfidfVectorizer_joblib.pkl')

sample_text = ["My cat only chews @apple cords. Such an #AppleSnob."]

prediction = loaded_model.predict(sample_text)
prediction
```

Out[58]: array(['3'], dtype=object)

Achievement of objective

- Insights gained can help business better understand customer sentiment and make Data driven decision to enhance products and service.

Conclusion

- Both Logistic Regression with grid search and Bagged Random Forest performed the best with Accuracy of 71%.
- The Logistic Regression with gridsearch achieved an accuracy of 0.713 and precision of 0.677 while Bagged Random Forest with accuracy of 0.708 and precision of 0.685
- There is still room for improvement, particularly in handling the positive class (5).

Recommendations

By doing the following,the model performance will improve:

- Experimenting with advanced feature engineering,
- Experiment further with different other models
- Implement further ensemble methods

Limitations

- Labeling tweets as "Positive", "Negative" or "No emotion" can be a highly subjective exercise. What I may think is a positive tweet, someone else may interpret as negative.
- The context of these tweets matter. Since we don't know the methodology of how the data was labeled, there could have been human error in labeling where a tweet that was intended to be sarcastic can be labeled incorrectly for example. This would negatively impact the quality of the data.

Next Steps

- Collecting additional data to improve the model.

Type *Markdown* and LaTeX: α^2